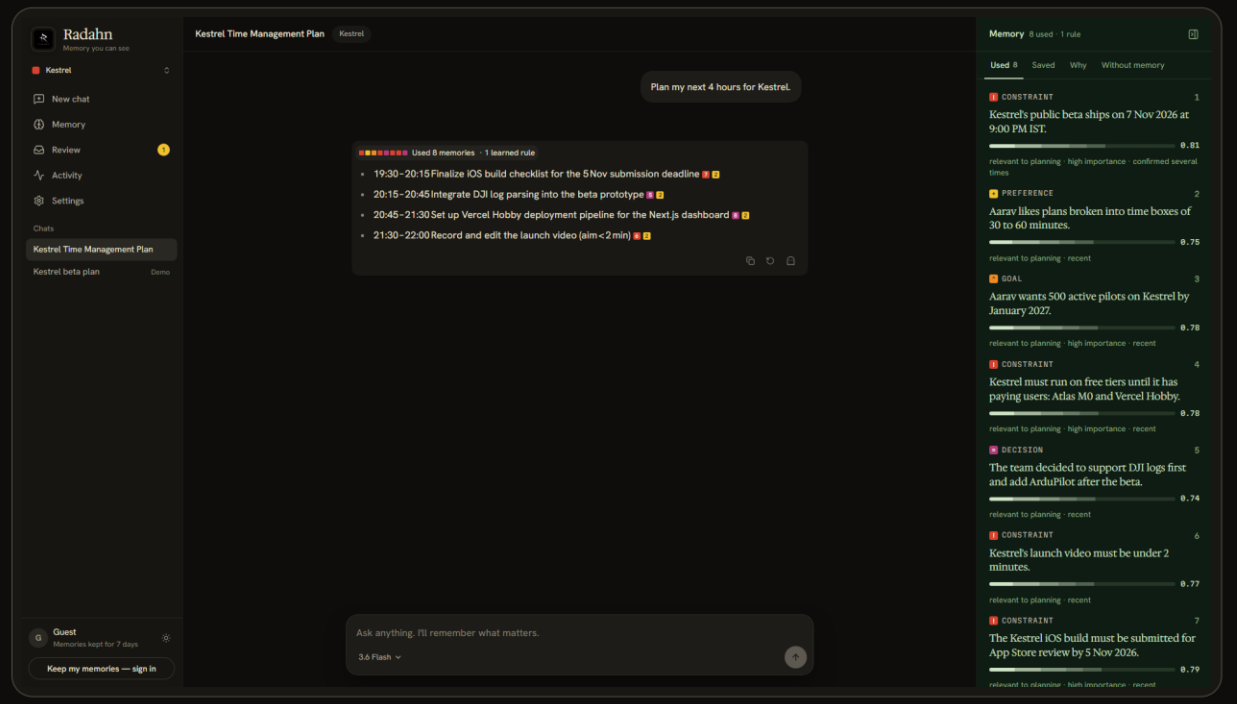
Radahn

An AI agent whose memory you can see, steer and trust.

radahn-ai.vercel.app

Try it as a guest. No sign-up.



Submitted to UNCHARTED 2026 · Agentic Memory track

Assistants remember you. Invisibly.

You can’t see what it kept

Memory is a hidden list of notes or summaries.

You can’t see why it changed an answer

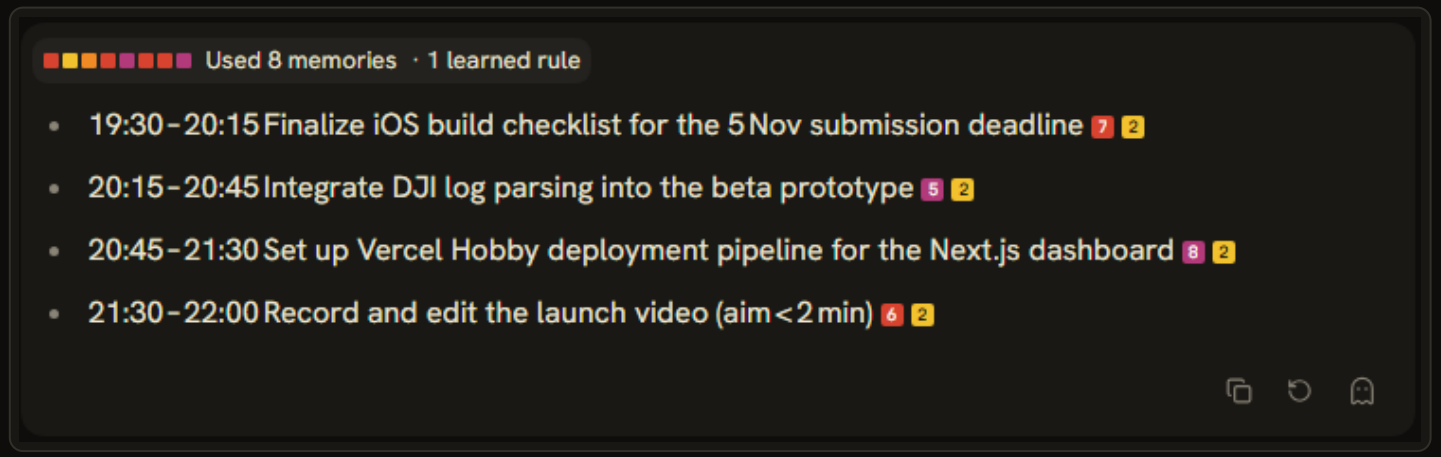
No receipts: which fact shaped which line, and how much.

You can’t fix one wrong fact

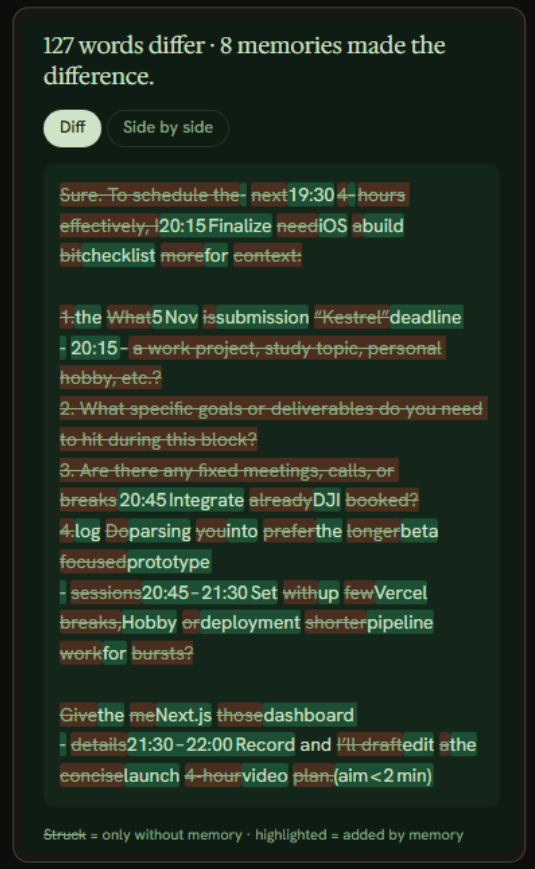
Old and new facts pile up; correcting one means starting over.

The problem isn’t forgetting. It’s memory you can’t inspect or steer.

Every answer shows its receipts



Coloured squares cite the exact memory behind each line.
One click asks again with memory off and highlights what changed.



WITHOUT MEMORY → DIFF

PRODUCT

What it does, in four verbs

01

Remembers

Small typed memories from every chat: preferences, decisions, deadlines, feedback.

02

Recalls and explains

Finds what matters for this message and shows the score behind every pick.

03

Corrects with history

A changed fact becomes history, never overwritten. Rewind to any earlier moment.

04

Forgets for real

Forget deletes. Expiry is automatic. Secrets are never stored.

WRITE PATH

How a memory is written

01

Privacy pre-scan

secrets removed before any model sees them

02

Extract

a small model proposes candidate memories

03

Policy gate

sensitive → Review; unsure → asks you

04

Reconcile

add · confirm · supersede · skip

05

Transaction

one current value per fact, history kept

06

Live event

change stream updates every screen

add confirm supersede skip

Runs after the reply, so it never slows you down. First memory lands ≈ 2 s later.

How a memory is found

Meaning
vector search

Keywords
Atlas Search

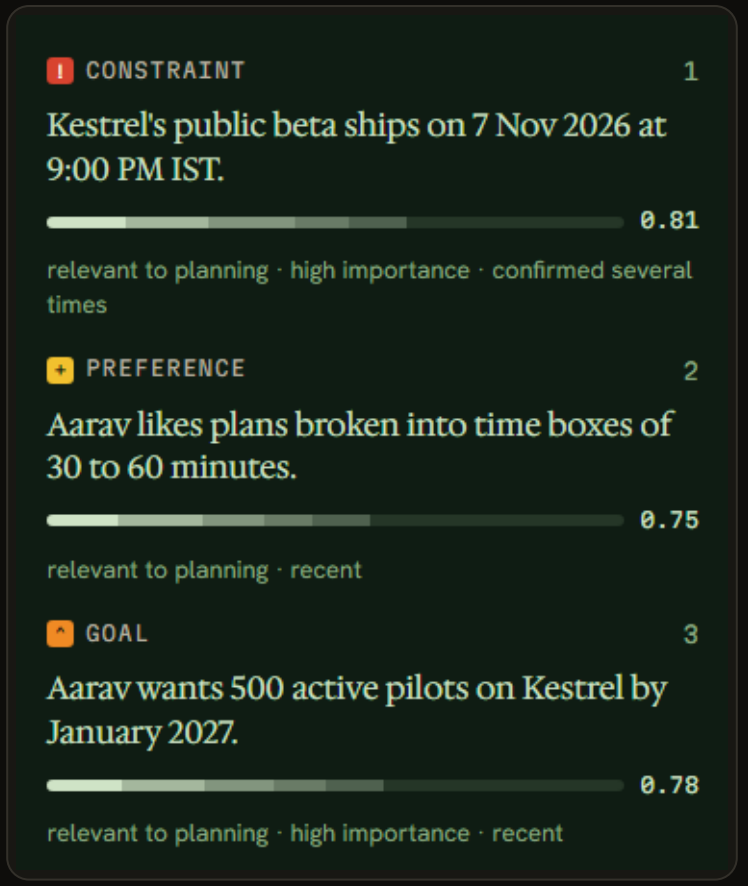
Time
only for past questions

Links
shared people and things

FUSE → SCORE → TOP 8

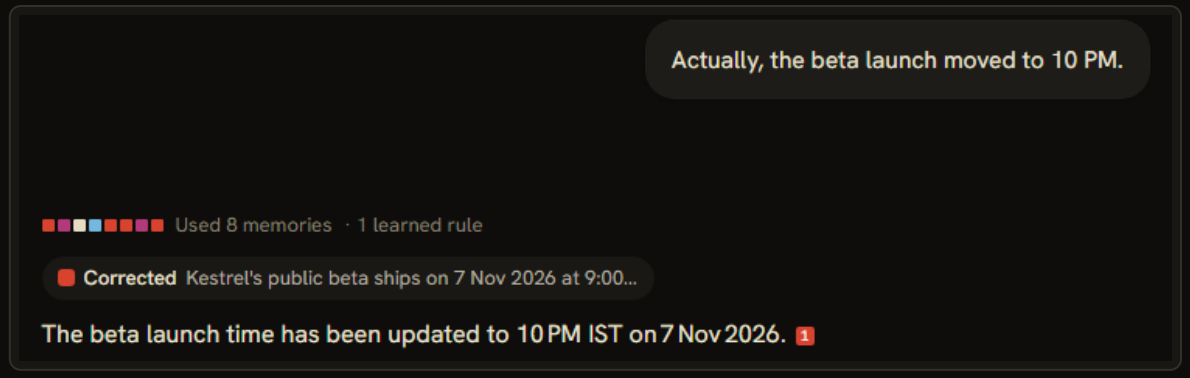
0.55 relevance
0.15 recency
0.15 importance
0.10 confidence

Weak matches dropped, near-duplicates removed.

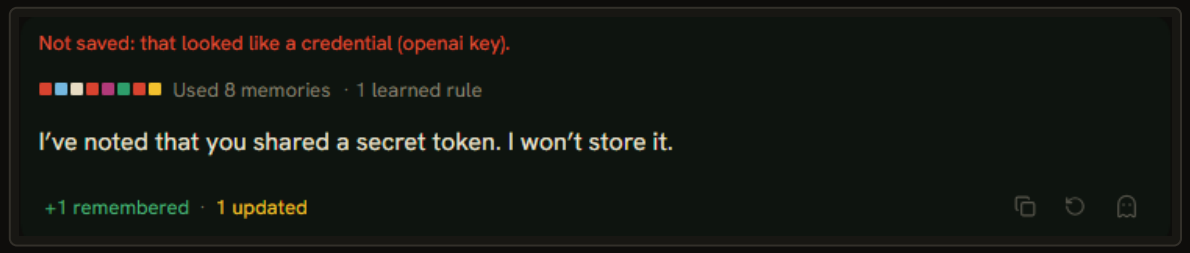


Every number is stored and shown in the Why tab, so every answer can be explained.

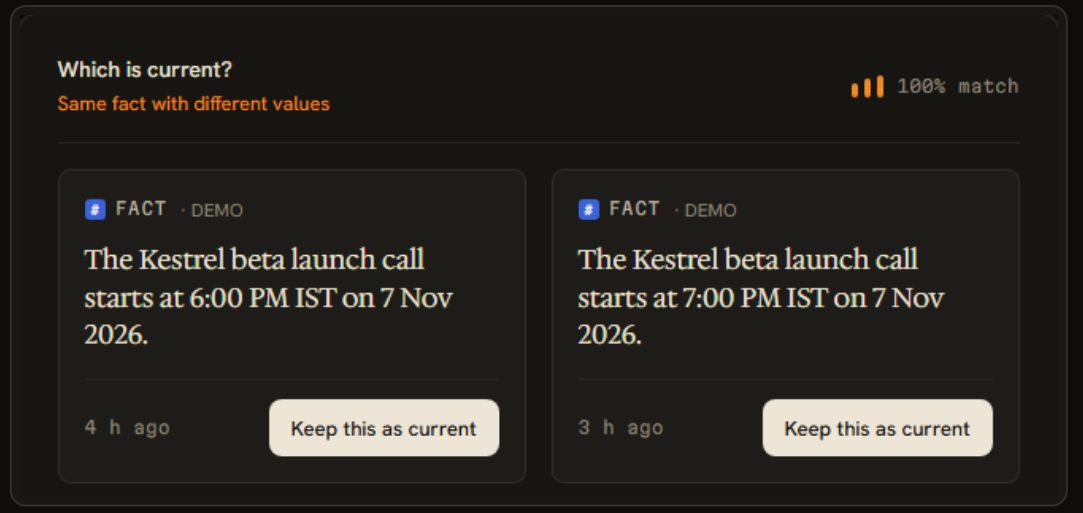
The hard cases, handled



CORRECTION → OLD VALUE KEPT AS HISTORY



SECRETS BLOCKED BEFORE THE MODEL SEES THEM



CONTRADICTIONS PUT TO YOU, NOT GUESSED

Expiry: stops at the set time; MongoDB TTL deletes it.
Forget: deleted for good; only a content-free event remains.

REFLECTION

It learns how you like answers

- “That study plan was too long to use.”
- “Just the bullets, please.”
- “A 12-step plan at 2 AM is overwhelming.”

FEEDBACK, REMEMBERED THREE TIMES

LEARNED RULE · WITH EVIDENCE

When Aarav asks for a plan, give at most 5 bullets, each with a time box.

Next plan: short, time-boxed, and the reply shows “1 learned rule”. The Without-memory diff proves the change.

ARCHITECTURE

Built on MongoDB Atlas

Vector Search

meaning, filtered by user and project

Atlas Search

exact names and codes

TTL indexes

automatic expiry

Transactions

corrections never half-done

Partial unique index

one current value per fact

Change streams

live UI and agents

Aggregations

Activity stats and insights

Schema validation

every memory well-formed

Next.js on Vercel → **MongoDB Atlas** → **Gemini + Groq models** · one database for it all

RESULTS

Measured on production

20/20	31/31	1.0	0	≈ 0.5 s
live memory tests	rule checks	recall@5	secrets stored	recall time

LIMITS

Free tiers only. Extraction is an AI and can err, so every save is visible and reversible.

NEXT

A reranking model, shared team memory, importing existing notes.

Try it: **radahn-ai.vercel.app** · Team: [Name 1] · [Name 2] · [Name 3]